

# Paxos Explained from Scratch

**Hein Meling** and Leander Jehl




University of  
Stavanger

*hein.meling@uis.no*

DAT520 Distributed Systems 2024

# Leslie Lamport

- ▶ Microsoft Research
- ▶ Many important contributions to distributed computing theory
- ▶ 2013 Turing Award winner
- ▶ But most know for  $\text{\LaTeX}$



## The Part-Time Parliament

LESLIE LAMPORT  
Digital Equipment Corporation

Recent archaeological discoveries on the island of Paxos reveal that the parliament functioned despite the peripatetic propensity of its part-time legislators. The legislators maintained consistent copies of the parliamentary record, despite their frequent forays from the chamber and the forgetfulness of their messengers. The Paxos parliament's protocol provides a new way of implementing the state-machine approach to the design of distributed systems.

Categories and Subject Descriptors: C2.4 [Computer-Communications Networks]: Distributed Systems—*Network operating systems*; D4.5 [Operating Systems]: Reliability—*Fault-tolerance*; J.1 [Administrative Data Processing]: Government

General Terms: Design, Reliability

Additional Key Words and Phrases: State machines, three-phase commit, voting

## Vertical Paxos and Primary-Backup Replication

Leslie Lamport, Dahlia Malkhi, Lidong Zhou  
Microsoft Research

9 February 2009  
corrected 26 August 2009

### The Paxos Register

Harry C. Li, Allen Clement, Amitanand S. Aiyer, and Lorenzo Alvisi  
The University of Texas at Austin  
Department of Computer Sciences  
{harry, aclement, anand, lorenzo}@cs.utexas.edu

## There Is More Consensus in Egalitarian Parliaments

Iulian Moraru, David G. Andersen, Michael Kaminsky  
Carnegie Mellon University and Intel Labs

### Cheap Paxos

Leslie Lamport and Mike Massa

MICROSOFT

### Fast Paxos

Leslie Lamport

## Paxos for System Builders

Jonathan Kirsch and Yair Amir

## Paxos Made Live - An Engineering Perspective

Tushar Chandra  
Robert Griesemer  
Joshua Redstone

June 20, 2007

## Paxos Made Moderately Complex Paxos Made Simple

Robbert van Renesse  
Cornell University  
rvr@cs.cornell.edu

March 25, 2011

Leslie Lamport

01 Nov 2001

## In Search of an Understandable Consensus Algorithm

Diego Ongaro and John Ousterhout  
Stanford University

(Draft of April 7, 2013, under submission to SOSP)

## When You Don't Trust Clients: Byzantine Proposer Fast Paxos

Hein Meling\*, Keith Marzullo<sup>†</sup>, and Alessandro Mei<sup>‡</sup>

\*Department of Electrical Engineering and Computer Science, University of Stavanger, Norway

<sup>†</sup>Department of Computer Science and Engineering, University of California San Diego, USA

<sup>‡</sup>Department of Computer Science, Sapienza University of Rome, Italy

Email: hein.meling@uis.no, marzullo@cs.ucsd.edu, mei@di.uniroma1.it

# What is Paxos and why is it Relevant?

- ▶ Fault tolerant consensus protocol

# What is Paxos and why is it Relevant?

- ▶ Fault tolerant consensus protocol
- ▶ Used to order client requests sent to a fault tolerant server
  - ▶ For example a fault tolerant resource manager

# What is Paxos and why is it Relevant?

- ▶ Fault tolerant consensus protocol
- ▶ Used to order client requests sent to a fault tolerant server
  - ▶ For example a fault tolerant resource manager
- ▶ Used in production systems: Chubby, ZooKeeper, and Spanner

# What is Paxos and why is it Relevant?

- ▶ Fault tolerant consensus protocol
- ▶ Used to order client requests sent to a fault tolerant server
  - ▶ For example a fault tolerant resource manager
- ▶ Used in production systems: Chubby, ZooKeeper, and Spanner
- ▶ It is always safe

# Objectives and Approach

- ▶ Explain Paxos
  - ▶ In a step-wise manner
  - ▶ With minimal changes in each step




# Objectives and Approach

- ▶ Explain Paxos
  - ▶ In a step-wise manner
  - ▶ With minimal changes in each step
- ▶ Objective
  - ▶ Understand why it works and why the solution is necessary
  - ▶ (not how to implement or formally prove it)


# Objectives and Approach

- ▶ Explain Paxos
  - ▶ In a step-wise manner
  - ▶ With minimal changes in each step
- ▶ Objective
  - ▶ Understand why it works and why the solution is necessary
  - ▶ (not how to implement or formally prove it)
- ▶ Approach
  - ▶ Starting from a simple client/server system
  - ▶ Build fault tolerant server (replicated state machine)
  - ▶ Construct Multi-Paxos
  - ▶ Decompose Multi-Paxos into Paxos


# A Stateful Service: *SingleServer*




# A Stateful Service: *SingleServer* (Subscript)



# A Stateful Service: *SingleServer* (Superscript)



## A Stateful Service: *SingleServer*



- ▶ Client  $C_2$  sees:  $\sigma^2$
- ▶ Client  $C_1$  sees:  $\sigma^{21}$ 
  - ▶  $\sigma^2$  is a prefix of  $\sigma^{21}$
- ▶ Corresponds to execution sequence:  $m_2m_1$


We Want to Make the Service Fault  
Tolerant!

# Assumptions

- ▶ Asynchronous environment: No bounds on
  - ▶ Processing delay
  - ▶ Communication delay
  - ▶ Clock drift
  - ▶ Unreliable failure detectors
- ▶ Unreliable communication
  - ▶ May take arbitrarily long to deliver msg
  - ▶ Msgs can be duplicated and lost
- ▶ Deterministic operations
- ▶ Processes may crash



# Fault Tolerance with Two Servers




- ▶ Client  $C_2$  sees:  $\sigma^2$
- ▶ Client  $C_1$  sees:  $\sigma^{21}$ 
  - ▶  $\sigma^2$  is a prefix of  $\sigma^{21}$
- ▶ Execution sequence:  $m_2m_1$

# Deterministic State Machine

- ▶ Service implemented as a deterministic state machine
- ▶ Processing requests yield unique state transitions:
  - ▶  $\sigma_1^2 = \sigma_2^2$  and  $\sigma_1^{21} = \sigma_2^{21}$ .
- ▶ Clients suppress duplicate replies


# Fault Tolerance with Two Servers: Whoops!



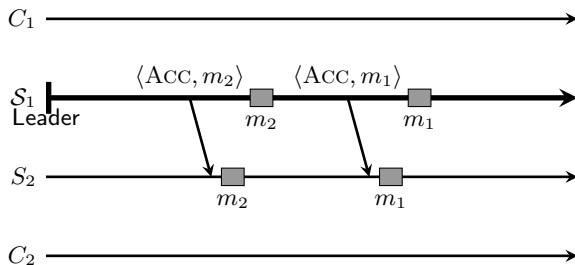
- ▶ Client  $C_2$  sees:  $\sigma^2 \sigma^{12}$ 
  - ▶  $\sigma^2$  not a prefix of  $\sigma^{12}$
- ▶ Client  $C_1$  sees:  $\sigma^1 \sigma^{21}$ 
  - ▶  $\sigma^1$  not a prefix of  $\sigma^{21}$
- ▶ Execution sequence
  - ▶  $S_1$ :  $m_1 m_2$
  - ▶  $S_2$ :  $m_2 m_1$

# We Need to Order Client Requests!


# Let's Designate a Leader to Order Requests




# Without Clients



# Problem: Also Accept Messages can be Reordered




# Add Sequence Numbers






# Discard Out-of-Order Messages



# Now with Clients



# Clients Observe The Same Server States as Before

- ▶ Client  $C_2$  sees:  $\sigma^2$
- ▶ Client  $C_1$  sees:  $\sigma^{21}$


# Clients Observe The Same Server States as Before

- ▶ Client  $C_2$  sees:  $\sigma^2$
- ▶ Client  $C_1$  sees:  $\sigma^{21}$
- ▶ However,  $S_2$  didn't execute  $m_1$ 
  - ▶ Q: What to do?

# Clients Observe The Same Server States as Before

- ▶ Client  $C_2$  sees:  $\sigma^2$
- ▶ Client  $C_1$  sees:  $\sigma^{21}$
- ▶ However,  $S_2$  didn't execute  $m_1$ 
  - ▶ Q: What to do?
  - ▶ A1: Buffer

# Problem: Message Loss – $S_2$ Won't Execute Anything




# Clients Observe The Same Server States as Before

- ▶ Client  $C_2$  sees:  $\sigma^2$
- ▶ Client  $C_1$  sees:  $\sigma^{21}$
- ▶ However,  $S_2$  didn't execute  $m_1$ 
  - ▶ Q: What to do?
  - ▶ A1: Buffer
  - ▶ A2: Retransmission mechanism

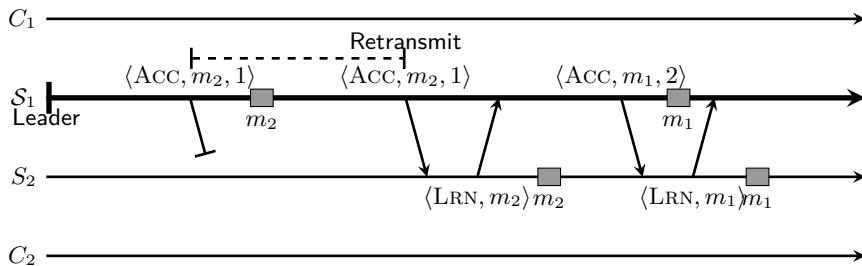
# We Need a Retransmission Mechanism!




# A Learn Stops Retransmission



# Don't Send New Accept Until Learn



# With Clients



# Recap

- ▶ A leader
  - ▶ Decides order of client requests
  - ▶ By sending an accept message to  $S_2$

# Recap

- ▶ A leader
  - ▶ Decides order of client requests
  - ▶ By sending an accept message to  $S_2$
- ▶ Sequence numbers
  - ▶ Cope with message reordering

# Recap

- ▶ A leader
  - ▶ Decides order of client requests
  - ▶ By sending an accept message to  $S_2$
- ▶ Sequence numbers
  - ▶ Cope with message reordering
- ▶ Retransmission mechanism
  - ▶ Cope with message loss
  - ▶ Leader only sends next accept when learn from  $S_2$

# Recap


- ▶ A leader
  - ▶ Decides order of client requests
  - ▶ By sending an accept message to  $S_2$
- ▶ Sequence numbers
  - ▶ Cope with message reordering
- ▶ Retransmission mechanism
  - ▶ Cope with message loss
  - ▶ Leader only sends next accept when learn from  $S_2$

Combination of mechanisms:  
*RetransAccept* protocol


# What About Server Crashes?




# Crash




# Crash: Leader Takeover




# Single Server Rule: Case 1




## Single Server Rule: Case 2




# Single Server Rule: Case 3



# Single Server Rule: Case 4 – A Problem



# Lock Service: Both clients get the lock




## Single Server Rule: Case 4 – A Problem

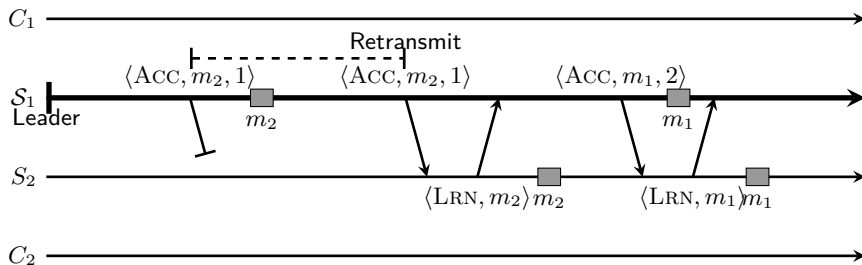
- ▶ Imagine that  $(S_1, S_2)$  is a fault tolerant lock service
- ▶ Both clients could have gotten the lock




# Solution: Leader Waits for Learn Before Executing




# Recall Earlier Version




# Now Leader Takeover is Safe




# Let's Add Client Messages



# Leader Remain in Control when $S_2$ Crash



# Let's Add Client Messages Again



# Recap: The Problem

- ▶ When we detect a server crash
  - ▶ Switch to *SingleServer* mode

# Recap: The Problem

- ▶ When we detect a server crash
  - ▶ Switch to *SingleServer* mode
- ▶ Problem with *RetransAccept* protocol:
  - ▶ Leader might have replied to a client and then crashed, without ensuring that  $S_2$  saw the accept



# Recap: The Problem

- ▶ When we detect a server crash
  - ▶ Switch to *SingleServer* mode
- ▶ Problem with *RetransAccept* protocol:
  - ▶ Leader might have replied to a client and then crashed, without ensuring that  $S_2$  saw the accept
  - ▶  $S_2$  takes over and may execute a different request in *SingleServer* mode

## Recap: WaitForLearn Protocol

- ▶ Leader waits for a learn from  $S_2$

## Recap: WaitForLearn Protocol

- ▶ Leader waits for a learn from  $S_2$
- ▶  $S_2$  can execute after seeing an accept from the leader
  - ▶ Because the accept is also an implicit learn

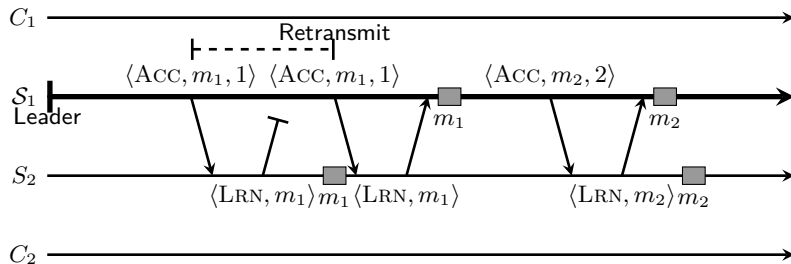
# Question

- ▶ Q: What happens if the learn message to the leader is lost?

# Question

- ▶ Q: What happens if the learn message to the leader is lost?
- ▶ A: The leader uses *RetransAccept*; the accept will be retransmitted. So no need for another retransmit protocol.

# If the Learn is Lost, Retrans will fix it



# Somewhat Rougher Road Ahead!

# False Detection

- ▶ So far we have assumed that failure detection is accurate



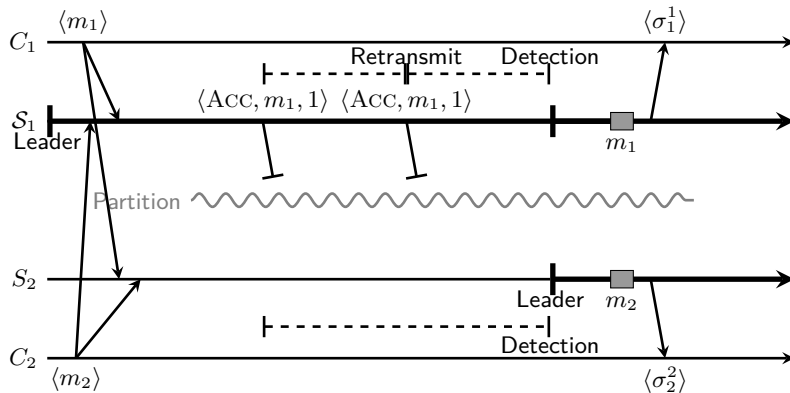
# False Detection

- ▶ So far we have assumed that failure detection is accurate
- ▶ But in an asynchronous environment
  - ▶ There is always a chance of false detection
  - ▶ Because it is impossible to pick the right timeout delay

# False Detection

- ▶ So far we have assumed that failure detection is accurate
- ▶ But in an asynchronous environment
  - ▶ There is always a chance of false detection
  - ▶ Because it is impossible to pick the right timeout delay
- ▶ We now consider false detection in the context of network partitions

# Problem: Network Partitions



# Network Partition

- ▶ Each server can switch to *SingleServer* mode (no coordination) and make progress

# Network Partition

- ▶ Each server can switch to *SingleServer* mode (no coordination) and make progress
- ▶ But it will lead to inconsistencies
  - ▶  $S_1$  has state  $\sigma^1$
  - ▶  $S_2$  has state  $\sigma^2$

# Network Partition

- ▶ Each server can switch to *SingleServer* mode (no coordination) and make progress
- ▶ But it will lead to inconsistencies
  - ▶  $S_1$  has state  $\sigma^1$
  - ▶  $S_2$  has state  $\sigma^2$
- ▶ Reconciling the state divergence
  - ▶ Involves rollback on multiple clients


# Network Partition

- ▶ Each server can switch to *SingleServer* mode (no coordination) and make progress
- ▶ But it will lead to inconsistencies
  - ▶  $S_1$  has state  $\sigma^1$
  - ▶  $S_2$  has state  $\sigma^2$
- ▶ Reconciling the state divergence
  - ▶ Involves rollback on multiple clients
  - ▶ Quickly becomes unmanageable

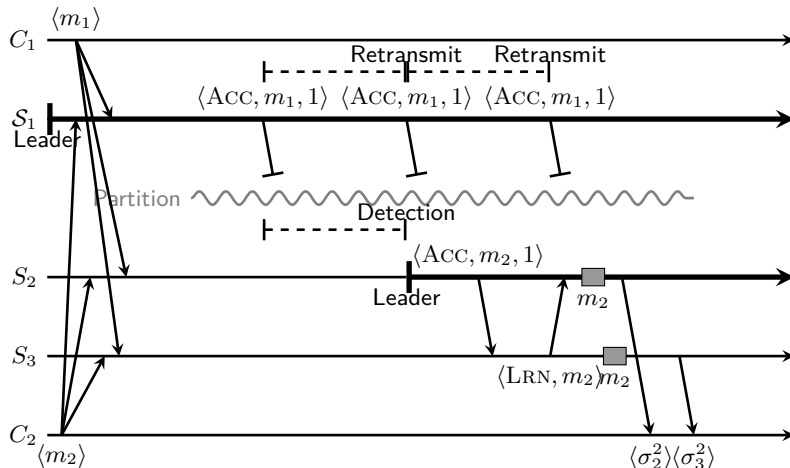
# We Want to Avoid Relying on Clients!




# Add Another Server; Make Progress in Majority Partition




# New Leader in Majority Partition



# WaitForLearn Without Partition



# WaitForLearn With Clients



## Recap: Network Partition

- ▶ We added another server,  $S_3$ 
  - ▶ To avoid rollback using clients

## Recap: Network Partition

- ▶ We added another server,  $S_3$ 
  - ▶ To avoid rollback using clients
- ▶ We still use the *WaitForLearn* protocol
  - ▶ To ensure that another server has seen the accept message

## Recap: Network Partition

- ▶ We added another server,  $S_3$ 
  - ▶ To avoid rollback using clients
- ▶ We still use the *WaitForLearn* protocol
  - ▶ To ensure that another server has seen the accept message
- ▶ Leader only needs to wait for *one* learn before executing the request
  - ▶ Allows the leader to make progress, when one of the other servers has crashed or is temporarily unavailable


## Recap: Network Partition

- ▶ We added another server,  $S_3$ 
  - ▶ To avoid rollback using clients
- ▶ We still use the *WaitForLearn* protocol
  - ▶ To ensure that another server has seen the accept message
- ▶ Leader only needs to wait for *one* learn before executing the request
  - ▶ Allows the leader to make progress, when one of the other servers has crashed or is temporarily unavailable
- ▶ But we still only tolerate one concurrent failure
  - ▶ Either a crash or a network partition

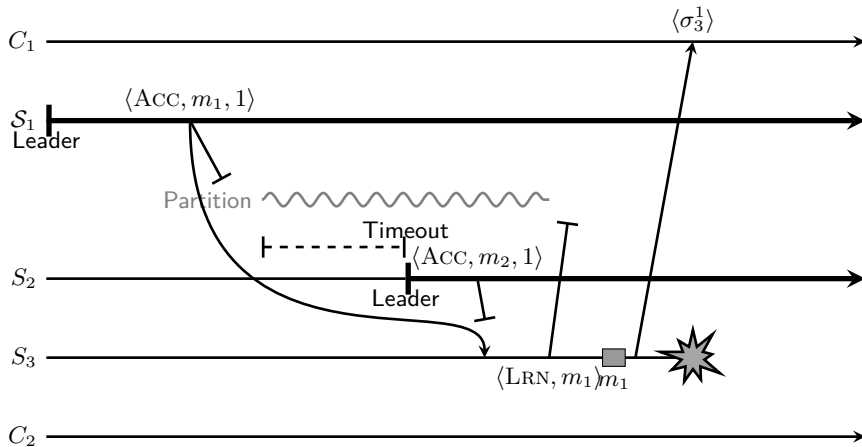


# What can go Wrong: Concurrent Crash and Partition

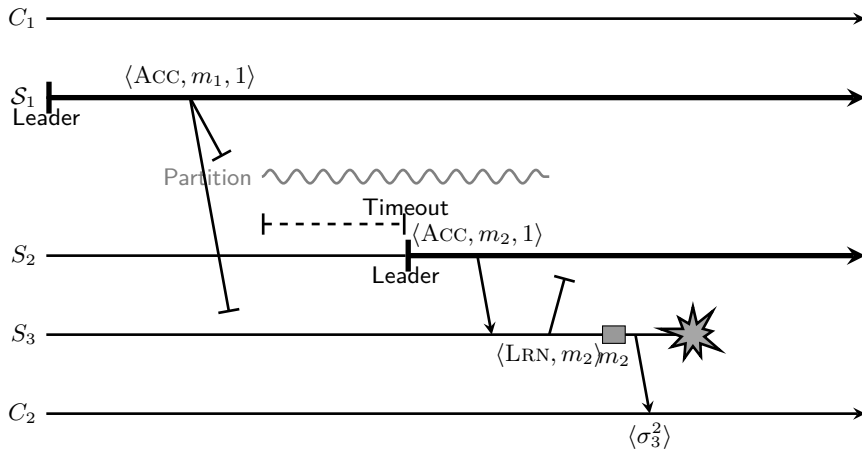
# Concurrent Crash and Partition



# Crash and Partition: Outcome 1 – $m_1$ Executed



## Crash and Partition: Outcome 2 – $m_2$ Executed



# Recap: Crash and Partition

- ▶  $S_3$  crashed
  - ▶ But *it could* have executed either  $m_1$  or  $m_2$
  - ▶ And replied to a client

## Recap: Crash and Partition

- ▶  $S_3$  crashed
  - ▶ But *it could* have executed either  $m_1$  or  $m_2$
  - ▶ And replied to a client
- ▶ Other servers cannot determine which message, if any, was executed

# Recap: Crash and Partition

- ▶  $S_3$  crashed
  - ▶ But *it could* have executed either  $m_1$  or  $m_2$
  - ▶ And replied to a client
- ▶ Other servers cannot determine which message, if any, was executed
  - ▶ Maybe we could talk to clients?
  - ▶ We don't want to rely on clients!

# Explicit Leader Change Mechanism

- ▶ Above problem is rooted in possibility of false detection
  - ▶ Can lead to several servers thinking they are leaders
  - ▶ And sending accept messages concurrently




# Explicit Leader Change Mechanism

- ▶ Above problem is rooted in possibility of false detection
  - ▶ Can lead to several servers thinking they are leaders
  - ▶ And sending accept messages concurrently
- ▶ It can be solved by an explicit leader takeover protocol


# Explicit Leader Change Mechanism

- ▶ Above problem is rooted in possibility of false detection
  - ▶ Can lead to several servers thinking they are leaders
  - ▶ And sending accept messages concurrently
- ▶ It can be solved by an explicit leader takeover protocol
- ▶ We need a way to
  - ▶ Distinguish messages from different leaders
  - ▶ Change the leader

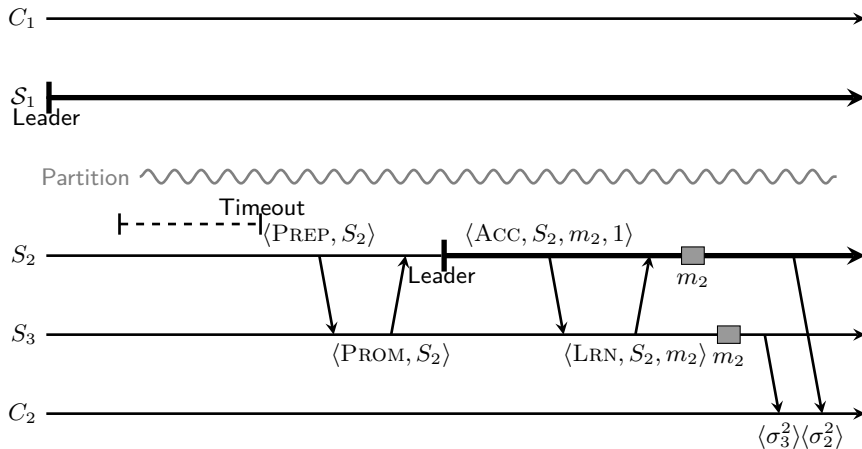
# Explicit Leader Change




# Leader Identifiers in Accept and Learn Messages




# With Client Replies




# Ignore Accept From Old Leader



# What Happens Now?




# Case 1: $S_3$ Takes Over?






## Case 2: $S_1$ Takes Over Again?



# Replace Leader Identifiers With Round Numbers



## Recap: Leader Change

- ▶ Added round number  $rnd$  in messages
  - ▶ To identify the leader
    - ▶  $\langle ACC, rnd, m, seqno \rangle$ : Sent by leader of round  $rnd$
    - ▶  $\langle LRN, rnd, m \rangle$ : Sent to leader of round  $rnd$

# Recap: Leader Change


- ▶ Added round number  $rnd$  in messages
  - ▶ To identify the leader
    - ▶  $\langle ACC, rnd, m, seqno \rangle$ : Sent by leader of round  $rnd$
    - ▶  $\langle LRN, rnd, m \rangle$ : Sent to leader of round  $rnd$
  - ▶ Round numbers are assigned:
    - ▶  $S_1$ : 1, 4, 7, ...
    - ▶  $S_2$ : 2, 5, 8, ...
    - ▶  $S_3$ : 3, 6, 9, ...
  - ▶ Skipping rounds is possible

# Recap: Leader Change

- ▶ Added round number  $rnd$  in messages
  - ▶ To identify the leader
    - ▶  $\langle ACC, rnd, m, seqno \rangle$ : Sent by leader of round  $rnd$
    - ▶  $\langle LRN, rnd, m \rangle$ : Sent to leader of round  $rnd$
  - ▶ Round numbers are assigned:
    - ▶  $S_1$ : 1, 4, 7, ...
    - ▶  $S_2$ : 2, 5, 8, ...
    - ▶  $S_3$ : 3, 6, 9, ...
  - ▶ Skipping rounds is possible
- ▶ Added two new messages
  - ▶  $\langle PREP, rnd \rangle$ : Request to become leader for round  $rnd$
  - ▶  $\langle PROM, rnd \rangle$ : Promise not to accept messages from a lower round than  $rnd$  (i.e. an older leader)

Let's Apply This Together  
With Accept and Learn

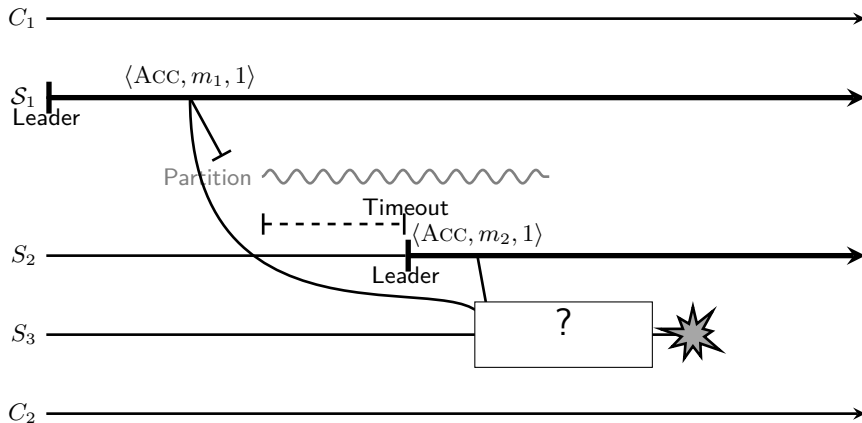
## $S_3$ Ignores Accept Message From Old Leader




Let's Recall the Problem we are Trying  
to Solve



# We Don't Know What $S_3$ Did Before Crashing




# Do We Know Now?




No we don't!

But it is Safe to Continue  
as If  $m_2$  Had Been Executed


# What Happens If $S_3$ Learn $m_1$ ?



# What Happens If $S_3$ Learn $m_1$ ?



# Does Leader Change Help?




No!

We Still don't Know What  
 $S_3$  Did Before Crashing.

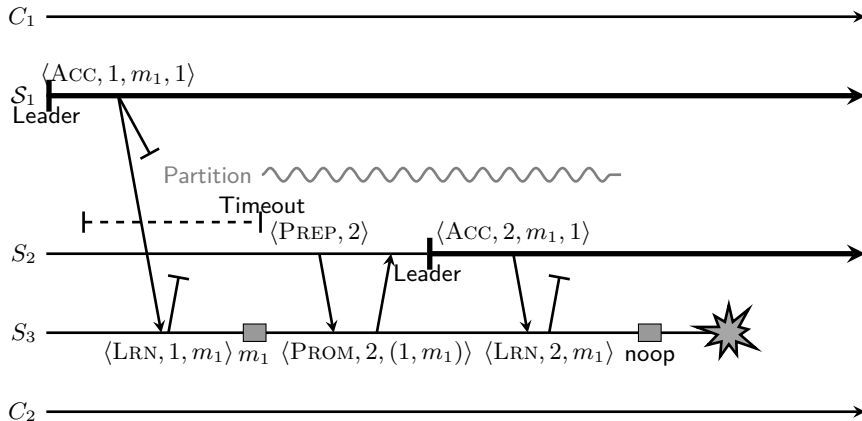


But the fix is Easy!

# Tell new Leader About Accepted Messages

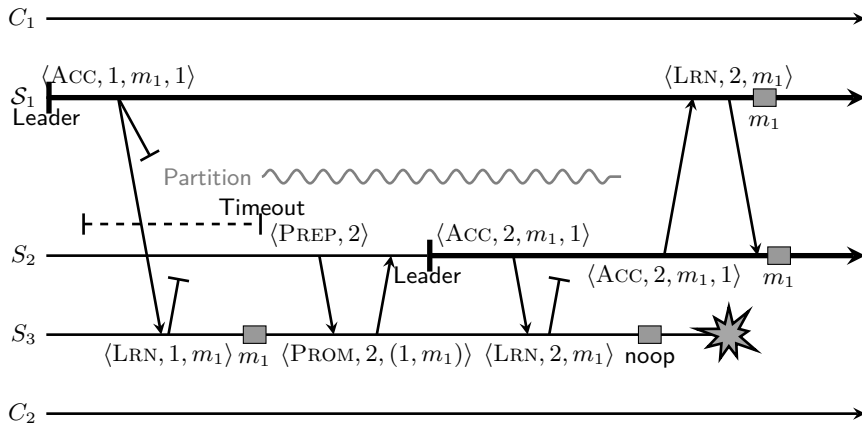


# The new Leader Resends Accept for Those Messages




Learn was Lost and  $S_3$  Crashed.  
Leader Still can't Execute  $m_1$ .

# Leader Also Resends Accept After Merge



# Promise from old Leader Includes Accepted Messages



## Recap: Leader Change 2

- ▶ Added information about accept from previous leader:  
 $\langle \text{PROM}, rnd, (1, m_1) \rangle$ 
  - ▶ Promise not to accept messages from a lower round than  $rnd$
  - ▶ Previous leader sent  $m_1$  in round 1
  - ▶ Typical naming:  $\langle \text{PROM}, rnd, (vrnd, vval) \rangle$

## Recap: Leader Change 2

- ▶ Added information about accept from previous leader:  
 $\langle \text{PROM}, rnd, (1, m_1) \rangle$ 
  - ▶ Promise not to accept messages from a lower round than  $rnd$
  - ▶ Previous leader sent  $m_1$  in round 1
  - ▶ Typical naming:  $\langle \text{PROM}, rnd, (vrnd, vval) \rangle$
- ▶ Leader resends accept for messages identified in the promise
  - ▶ After receiving the promise
  - ▶ After a partition merge




# What About More Than one Crash?


# What About More Than one Crash?

- ▶ Increase the number of servers
- ▶ Limit progress to a majority partition:
  - ▶ Can only tolerate that fewer than half of the servers fail
  - ▶ To tolerate  $f$  crashes, we need  $2f + 1$
  - ▶ Majority:  $f + 1$

# With Five Servers



# With Five Servers, $S_2$ Cannot Execute After Accept



# With Five Servers, $S_2$ Cannot Execute After Accept

- ▶ A combination of message loss and crashes
  - ▶ Prevent non-leader servers from executing after receiving an accept


# With Five Servers, $S_2$ Cannot Execute After Accept

- ▶ A combination of message loss and crashes
  - ▶ Prevent non-leader servers from executing after receiving an accept
  - ▶ This was not necessary for the three server case
    - ▶ The accept from the leader is an implicit learn
    - ▶ And together with its own "learn", can execute!

# With Five Servers, $S_2$ Cannot Execute After Accept


- ▶ A combination of message loss and crashes
  - ▶ Prevent non-leader servers from executing after receiving an accept
  - ▶ This was not necessary for the three server case
    - ▶ The accept from the leader is an implicit learn
    - ▶ And together with its own "learn", can execute!
- ▶ There are two solutions:
  - ▶ Wait for all-to-all learn
  - ▶ Wait for commit from leader

# All-to-All Learn Before Execute






# Await Commit Before Execute




# Wrapping it up!


# Multi-Paxos




# Multi-Paxos with Clients




# Paxos



# Paxos with Failure



# Paxos



That's It! Thank You!